



VECTOR
INSTITUTE

Bias in AI:

Week #5: Course Review & Capstone Project

Instructor:

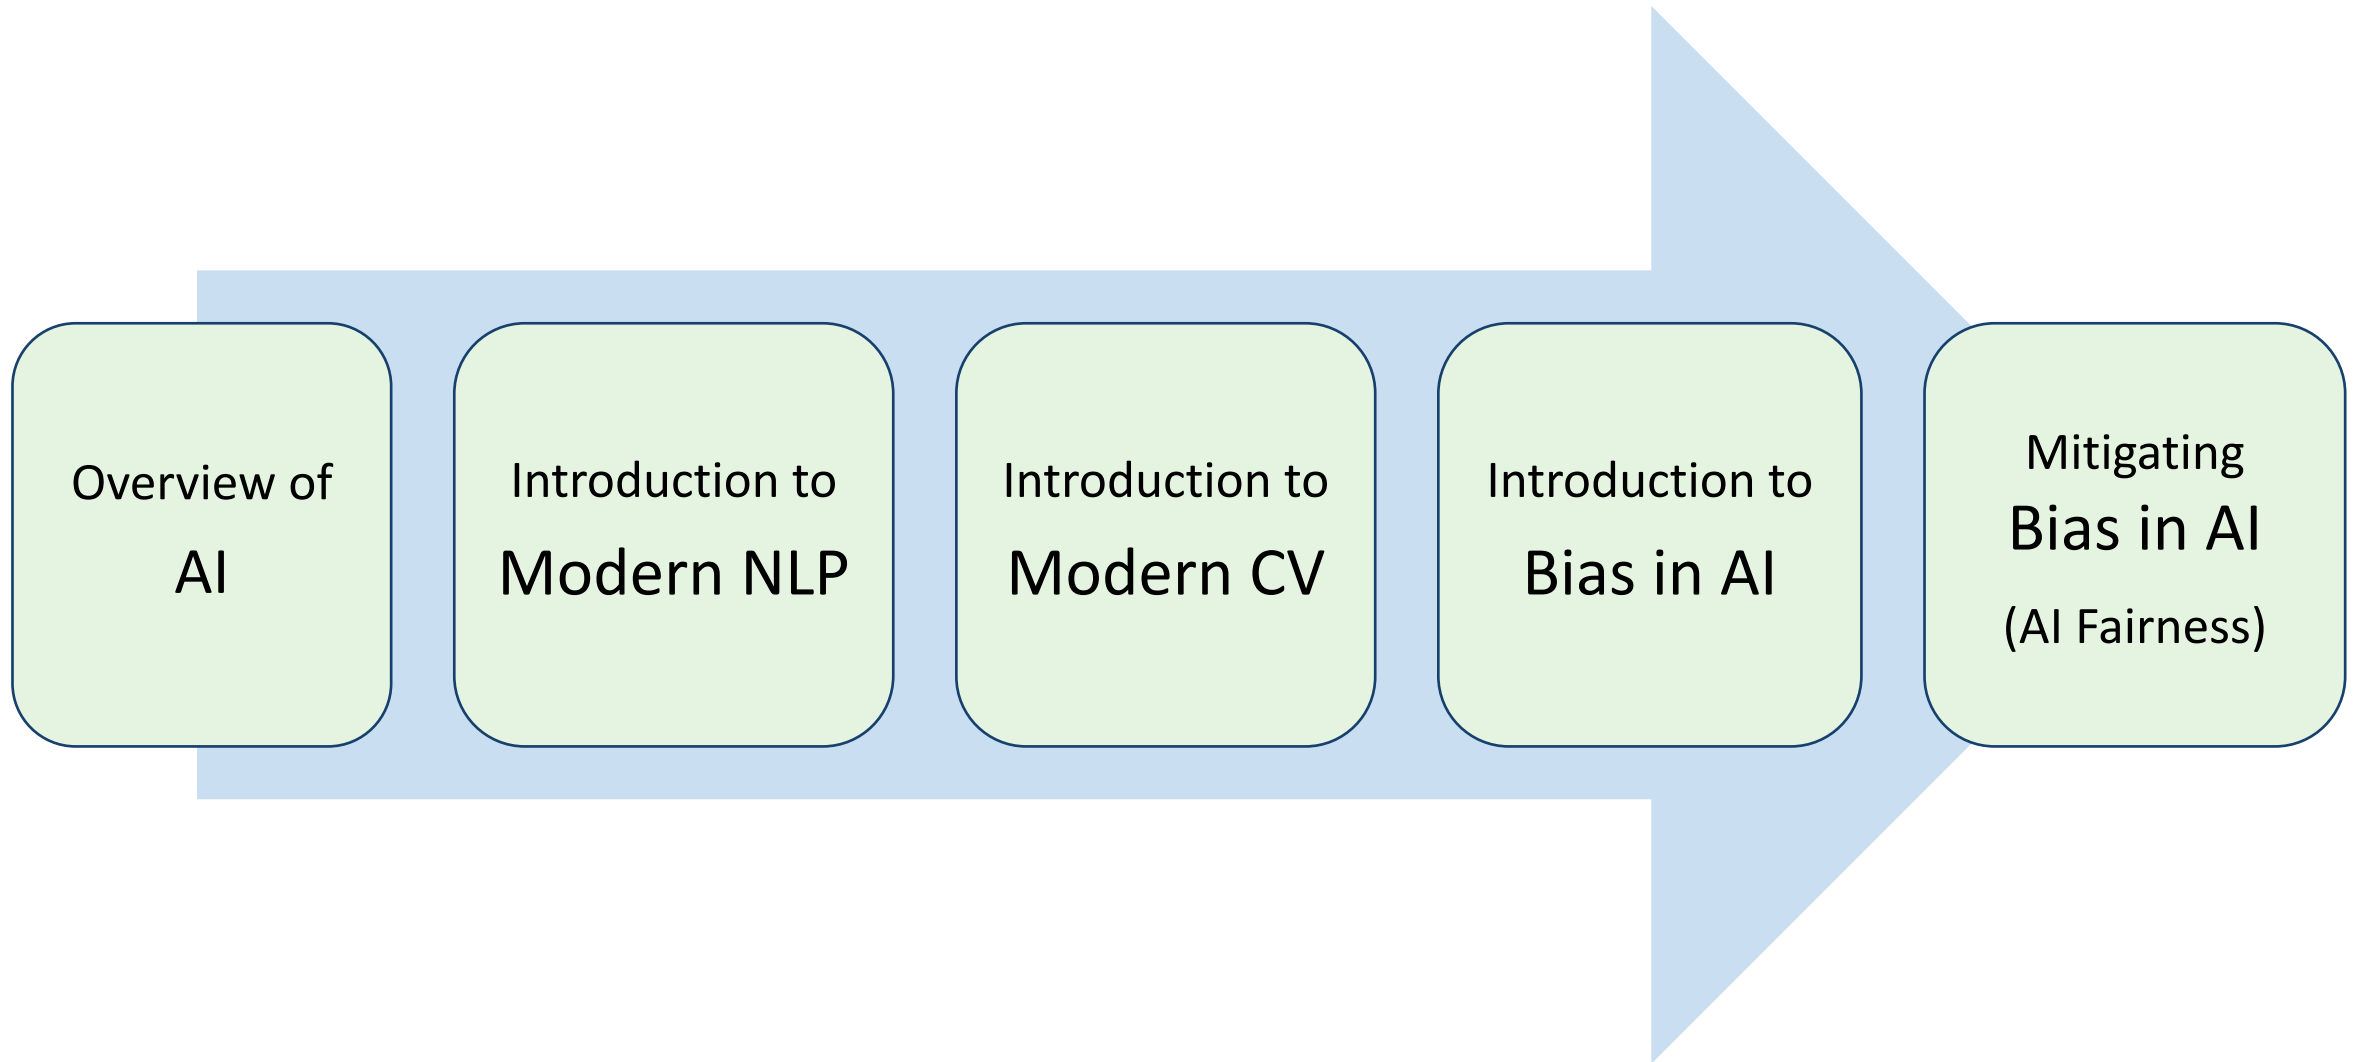
Sayyed Nezhadi

Winter 2022

What was the Goal?

- Create **awareness** about Bias in AI, sources and types, and potential harmful impacts
- Introduce potential approaches to identify and mitigate bias in AI systems
- Practical Examples in NLP & Computer Vision
- **Apply** learnings to your own businesses, to innovate, improve processes, solve problems

What We Discussed (Review)



What We Discussed (Review)

Overview of
AI

Introduction to
Modern NLP

- NLP Applications
- Word and Text Representation
- Word Embedding
- Text Classification
- Sequential Models (Recurrent Neural Networks)
- Language Model
- Encoder-Decoder Models (Seq2seq)
- Using Attention Methods
- Transformers (All you need is attention)
- Pre-trained Language Models (BERT and GPT-n)
- Assignment: Building a Chatbot using GPT-2 and analyzing the Bias in the Results

What We Discussed (Review)

- CV Applications
- Image Representation
- Image Classification
- Convolutional Neural Networks
- ImageNet Models
- Transfer Learning and Fine-tuning
- Object Detection Methods
- Image to Text (Image Captioning)
- Encoder-decoder Models
- Attention Methods
- Generative Models
- Generative Adversarial Networks

Introduction to
Modern CV

- Assignment:
 - Training a computer vision model for gender classification
 - Analyzing the bias in the results

What We Discussed (Review)

- Examples of High-profile Cases that AI Failed due to Harmful Bias
- Types of Bias Harms
- Good Bias vs. Bad Bias
- The Bias in the Decision Cycle
- Network Effect and Examples
- Biases in Data (e.g., Reporting Bias, Selection Bias)
- Examples of Bias in Language Models
- Examples of Bias in Image Representation
- Algorithmic Bias
- Bias in Interpretation (e.g., Automation Bias, Confirmation Bias, Overgeneralization, Correlation vs. Causation)
- Assignment: a report on how bias affects your organization and industry (technical and policy mitigation strategies)

Introduction to
Bias in AI

Mitigating
Bias in AI
(AI Fairness)

What We Discussed (Review)

-
- Discussion: Why Should We Care?
 - Regulated domains and classes
 - Team's Diversity
 - Challenges (Fairness Definition, Governance, Profit Trade-off)
 - End-to-End Bias Avoidance/Mitigation
 - Ethically Aligned Design
 - Documenting Datasets
 - Data Privacy
 - Model Transparency and Explainable AI
 - Fairness Criteria, Limitations, and Impossibility Theorem
 - Likelihood vs. Real Outcome
 - Equality vs. Equity
 - Individual vs. Group Fairness
 - Tools to Evaluate and Visualize Fairness
 - General Methods to Satisfy Fairness
 - Examples of Fair NLP
 - Examples of Fair Computer Vision

Overview of AI Fairness

Mitigating
Bias in AI
(AI Fairness)

Final Thoughts

Class Discussion

Capstone Project (one project by company*)

- Subject: Identify and reduce bias in your AI models.
- Objectives:
 - Explore and identify a possible bias issue in your company's AI models, products, or processes
 - Assess the type and the source of the bias, the potential harmful effects, the business and societal impacts, and associated regulations
 - Explore the potential technical and policy mitigation strategies
 - Choose the best feasible strategies to make the model/product/process more fair
 - Apply the selected technical strategies
 - Analyze the results and their impacts and/or trade-offs
 - Define future direction

Project Report

- Your project report should be written like a research-based business paper and should include the following sections :
 - Executive Summary (~250 words)
 - Introduction
 - Problem Definition
 - Model, Product, or Process
 - Explaining Data (e.g., Training Data)
 - Identified Bias
 - Impact Analysis
 - Mitigation Strategies
 - Experiments
 - Results
 - Conclusion & Future Work
 - References

Project Report

- Due by the end of Friday Apr 8th, 2022
- 8-10 pages long (excluding references)
- No need to include super sensitive content/ trade secrets
- The report should be created by LaTeX (if you don't have LaTeX on your computer, use <https://www.overleaf.com/>)
- Follow the suggested format by NeurIPS conference
- You can find the style files for NeurIPS 2020 here:
<https://nips.cc/Conferences/2020/PaperInformation/StyleFiles>

Final Presentations

- Presentation Date: Tuesday April 12th, 2022
- Presentation Length: 10 minutes followed by 5 minutes Q&A
- Capstone Panel: Vector Institute & IRAP Stakeholders
- Submission Deadline: Presentation documents can be submitted after the presentations by the end of Friday April 8th, 2022

Evaluation (Rubric)

Item	L4: Excellent (100%)	L3: Good (75%)	L2: Satisfactory (50%)	L1: Needs Improvement (25%)
Identifying Bias (15%)	Clearly defined problem statement, evaluated multiple candidates, selected an impactful case	Good problem statement, evaluated a couple of candidates, selected a moderately impactful case	Loosely defined problem statement, evaluated only one use-case, selected a case with low impact	No problem statement, No exploration / evaluation, selected a case with low impact
Impact Analysis (15%)	Clearly identified the type and source of the bias, analyzed the harms, impact analysis using an impact assessment framework	Identified the type and source of the bias, good impact analysis but without a standard framework	Either type or source of the bias is identified, loosely impact analysis	Type and source are not defined, no impact analysis is done, impacts are only named
Mitigation Strategy (15%)	Explored the potential technical and policy mitigation strategies, reasoned the strategy selection clearly with limitations and strengths	Explored the potential technical and policy mitigation strategies, reasoned the strategy selection	Explored either the potential technical and policy mitigation strategies, loosely reasoned the strategy selection	Selected the strategy with no exploration and without reasoning

Evaluation (Rubric) - continued

Item	L4: Excellent (100%)	L3: Good (75%)	L2: Satisfactory (50%)	L1: Needs Improvement (25%)
Analyzing the Results (15%)	Applied the strategy completely, Analyzed the results clearly, evaluated impacts and trade-offs, drew concrete conclusions	Applied the strategy, Analyzed the results, reviewed some impacts and trade-offs, drew conclusions	Applied the strategy loosely, Analyzed the results, drew conclusions	Incomplete application, provided the results without analysis and without impact analysis, conclusions without strong evidence
Project Report (20%)	Well formatted, included executive summary and all other listed sections (including references) within the page limit, clearly articulated, used proper visualization to convey the message, submitted on time	Well formatted, included executive summary and most other listed sections within the page limit, used some visualization to convey the message, submitted on time	Included executive summary and most other listed sections, used some visualization to convey the message, submitted on time	Formatted incorrectly, no executive summary or conclusion, no visualization to convey the message, submitted late
Presentation (20%)	Well formatted, used effective visualizations, conveyed the message clearly, presented within the time limit	Well formatted, used some visualizations, conveyed the message, presented within the time limit	Moderately formatted, used some visualizations, conveyed the message, presented close to the time limit	Not formatted properly, Mostly text, difficult to understand the message, too short or long